

SE350 – Project Overview

Third Part

Fall 2012

Outline

1. Processor Management
2. Scheduling
3. Initialization
4. Memory Management
- 5. Inter Process Communication**
- 6. Interrupts Handling**
- 7. Timing Services**

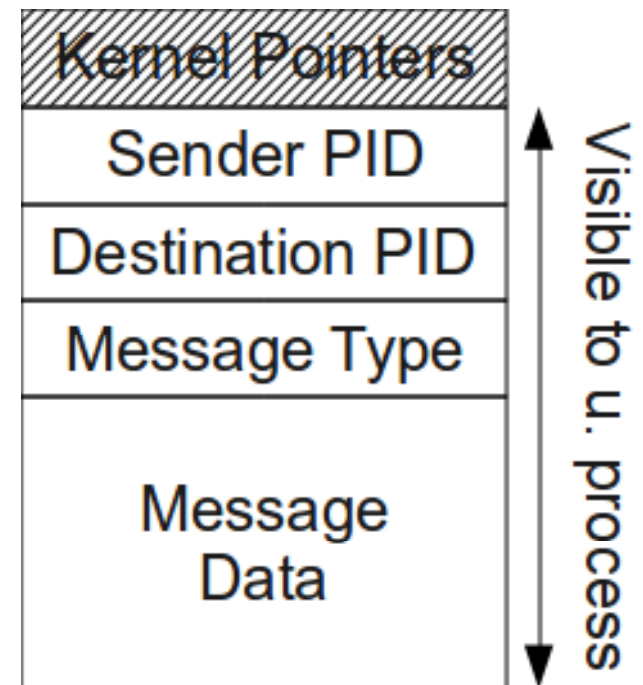
5. Inter Process Communication (IPC)

Requirements:

- Message-based, asynchronous IPC
- Messages are carried in “envelopes”

Procedure:

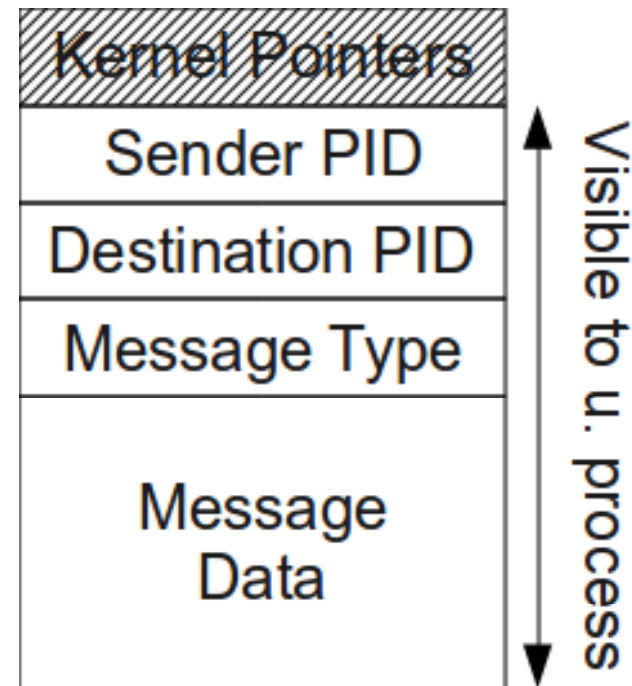
- Process writes into envelope
- Process invokes API function
`send_message(proc_id, envelope)`
- Other process invokes API function
`env = receive_message()` and blocks if no message is available



5. IPC: Message Envelopes

Message envelopes are managed by kernel:

- Create appropriate number of envelopes during init.
- Process allocates envelope using public API function `request_memory_block()`
- Process deallocates envelope using public API function `release_memory_block(env)` when not longer needed



5. IPC: Exchanging Messages

Requirement: messaged-based, asynchronous IPC

Design decisions:

- Non-blocking send
- Blocking receive

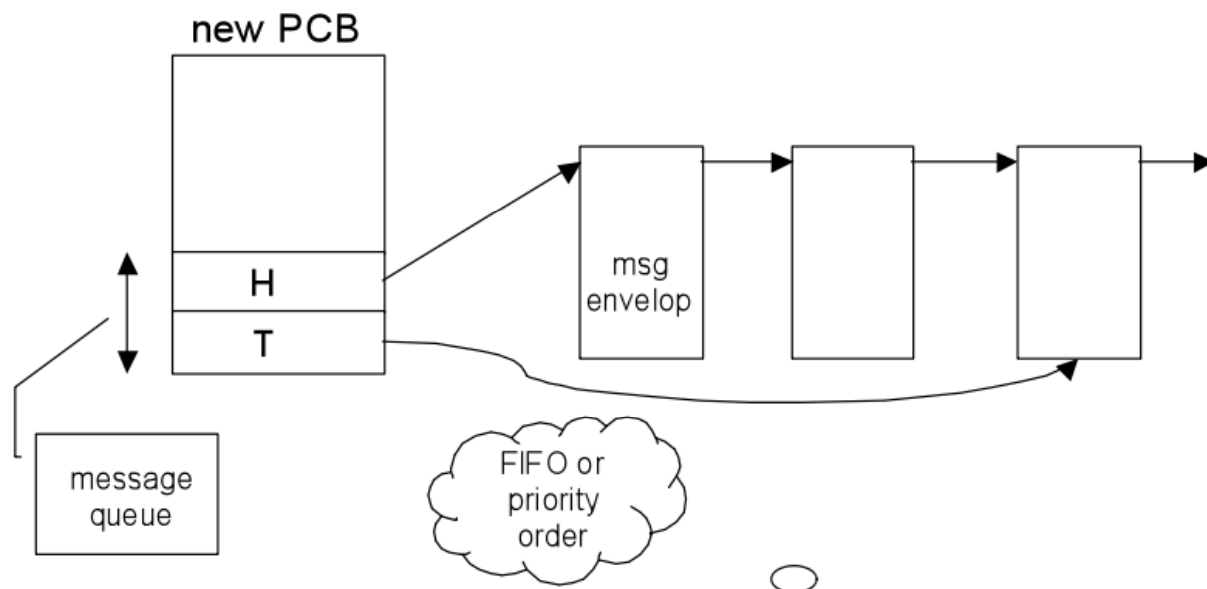
Issues:

- How do we handle multiple messages sent to a process?
- How does kernel track such messages?

5. IPC: Waiting Messages for Process

Idea: have each process have a queue of waiting messages

- **Extend PCB as follows**



5. IPC: send_message(pid, env)

```
send_message( target_pid, env) : {  
    atomic(on);  
    set sender_procid, destination_procid..  
    ...fields in env  
    target_proc -> convert target_pid ...  
    ...to process obj/PCB ref  
    enqueue env onto the msg_queue of target_proc  
  
    if ( target_proc.state is blocked_on_receive) {  
        set target_proc state to ready  
        rpq_enqueue( target_proc );  
    }  
    atomic(off);  
}
```

- Send never blocks!
- Send performs check if target_pid is blocked, if yes, it is unblocked and put into the ready state

5. IPC: receive_message()

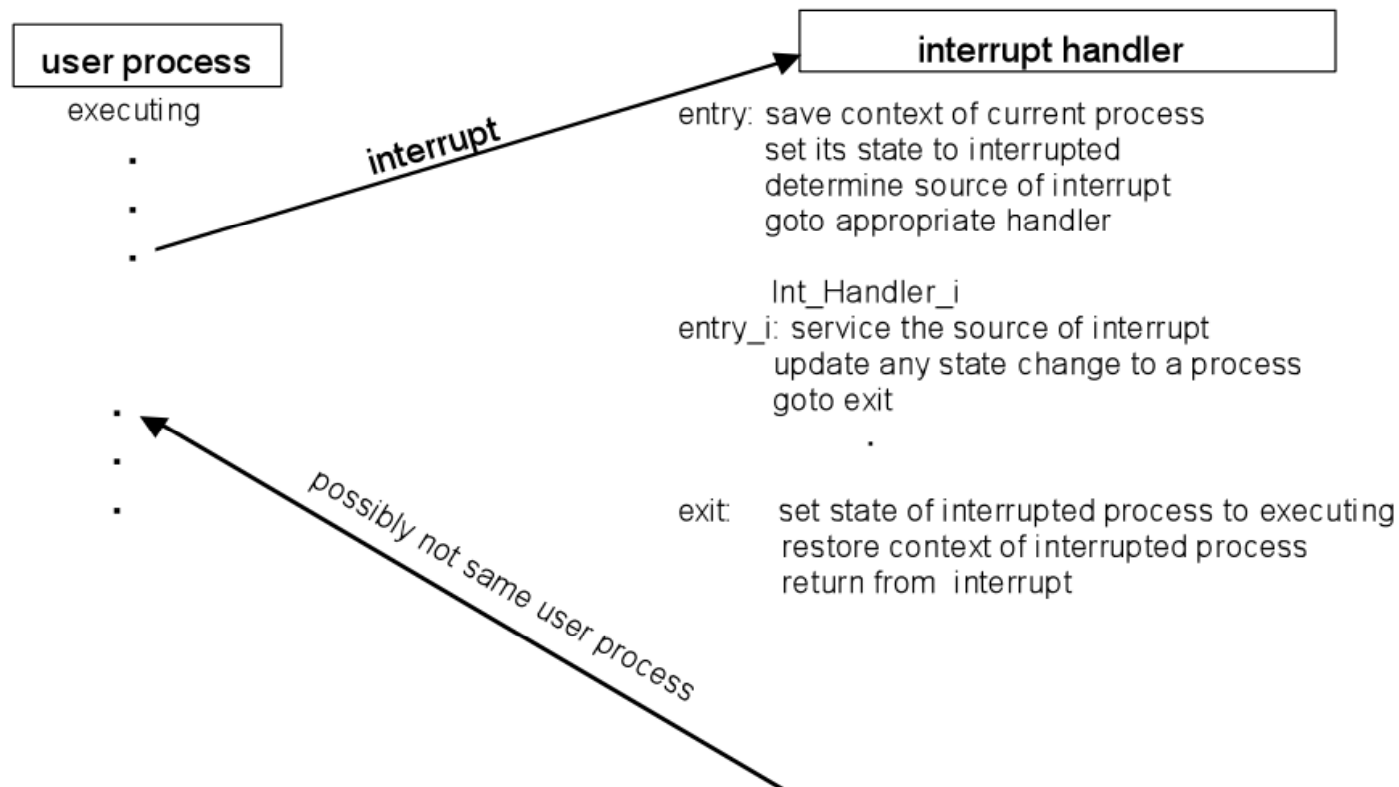
```
*env receive_message() {  
    atomic(on);  
    while ( current_process's msg_queue is empty) {  
        set state of current_process to blocked_on_receive  
        process_switch( );  
  
        *** return here when this process executes again  
    }  
  
    env -> dequeued envelope from the process' message queue  
    atomic(off);  
    return env;  
}
```

- When no messages available, process will block!
- Issue: Where should the blocked processes go?

6. Interrupt Handling & Atomicity

Interrupt = HW message, usually requiring short latency and quick service

- Interrupts invoke pre-registered procedures that “interrupt” currently executing code



6. Interrupt Handler: Design

Requirement: Interrupt handler must interact with RTX.

Design: i-process


- i-process is scheduled by interrupt handler
 - Is always ready to run, but not in ready queue
- i-process never blocks when invoking a kernel primitive
 - Primitives which can block the i-process need to be adapted (e.g., IPC)
- The interrupt handler starts the appropriate i-process
- Each i-process has a PCB

6. Interrupt Handler: Example

```
interrupt_handler {  
  
    set the state of current process to interrupted  
    save_proc = current_process  
  
    select interrupt source  
    A:  
        context_switch (i_proc_A)  
        break  
    ...  
  
    Z:  
        context_switch (i_proc_Z);  
        break  
    end select  
  
    //code to save context of interrupt handler (i_process)  
    current_process = save_proc;  
    context_switch (save_proc);  
    //perform a return from exception sequence  
    //this restarts the original process before i_handler  
}
```

7. Timing Service: Design

- Timing service is fundamental in RTX
- Timing service i-process receives message with service request (i.e., `delayed_send(PID, env, delay)` API), which is non blocking!!!
- After expiration time the timing i-process sends original message to the destination
- Timeout service maintains requests in sorted list:



7. Timing Service: i-process

At each clock tick (i.e. interrupt), the timeout process:

- Increments `current_time`
- Invokes `receive()` to see new requests (non-blocking!)
- If any new requests, it adds them to the list
- Checks if any timing requests have been expired
 - If yes, send the message to the destination.

7. Timing I-Process Example

```
timeout_i_process:
{
    env = receive(); //to get pending requests
    while (env is not null)
    {
        //code to insert the envelope into the timeout_list

        env = receive(); //see if any more requests
    }

    if (timeout_list is not empty)
    {
        while (head_of_timeout_list.expiry_time <= cur_time)
        {
            env = timeout_list.dequeue();
            target_pid = env.destination_pid;
            send( target_pid, env ); //forward msg to destination
        }
    }
}
```

