

**Department of Electrical and Computer Engineering
University of Waterloo**

SE 350 Project Description

1 Introduction

In this project, you will design a small real-time executive (RTX) and implement it on an MCF5307 based microcontroller board. The executive will provide a basic multiprogramming environment, with 5 priority levels, preemption, simple memory management, message-based interprocess communication, a basic timing service, system console I/O and debugging support.

Such an RTX is suitable for embedded computers which operate in real time. A cooperative, non-malicious software environment is assumed. The design of the RTX should allow its placement in ROM. Applications and non-kernel RTX processes must execute in the “user” mode of MCF5307. The RTX kernel will execute in the “supervisor” mode.

The description of the microcomputer and other project related information is available on the course web site (<http://www.ece.uwaterloo.ca/~ece354>) Project section. It has one megabyte of RAM for use by the RTX and application processes. It contains a fully integrated timer, a dual UART and several other peripheral interface devices. The board has two RS-232 interfaces, one for the system console (see later) and the other for software development/debugging support.

2 Summary of RTX Requirements

2.1 Scheduling Strategy

Four user priority levels plus an additional “hidden” priority level for the Null process, preemption, no time slicing, FIFO discipline at each priority level.

2.2 RTX Primitives and Services

Refer to the section RTX Primitives and Services.

2.3 RTX Footprint and Processor Loading

A reasonably 'lean' implementation is expected.

2.4 Error Detection and Recovery

At minimum, the RTX kernel must detect one type of error: an attempt to *send_message* to or *set_process_priority* of a non-existent process_id. The primitive will return an error code (a non-zero integer value). No error recovery is required. It may be assumed that the application processes can deal with this situation.

3 Organization and Deliverables

3.1 Project Groups

The project is done in groups of at most four. A group of three members is acceptable; however, there is no reduction in project deliverables. A group of two members or one member is not recommended.

3.2 Deliverables

The project has seven deliverables:

1. RTX Project P1 Implementation: The source code for serial port, timer and process switch.
2. RTX Project Software Design Description: This document describes the RTX design. It should include:
 - a structural description of the design (procedures and their interconnect; data structures; processes);
 - a functional description of all procedures (pseudo code; show all input/output parameters and globals);
 - implementation, testing and measurement plan (include responsibilities of individual team members).

It is important to keep this document reasonably small (no more than 30 pages, not including appendices). A design description should be shorter than the actual implementation.

3. RTX Project P2 Implementation: This is the source code of a simplified version of the final RTX. Only limited features are required to be implemented. The detailed requirements of the simplified RTX implementation will be given to you before or during the third week of the term.
4. RTX Project P2 Demonstration: The operation of the simplified RTX on the microcontroller will be demonstrated to the TAs in LAB. All group members must attend the demonstration.
5. RTX Project P3 Implementation: This is the final preemptive RTX source code to implement the specifications in sections 4, 5 6, and 7 based on the simplified RTX implementation done previously. It is possible that you will discover gaps or flaws in your original design during the implementation phase. You are allowed to adjust your implementation to reflect such discoveries. It is expected that you will document any major changes in the Final Report.
6. RTX Project P3 Demonstration: The operation of the final preemptive RTX on the microcontroller must be demonstrated to the TAs and/or to an instructor. At minimum, this will include the demonstration of continuous operation of the set of demo/test processes described in a later section. All group members must attend the demonstration.
7. Final Project Report: This report will include
 - a) An overview of major design changes made during implementation, if any.
 - b) The measured times for each of the *send_message*, *receive_message* and *request_memory_block* primitives, and a check of the reasonableness of the values measured
 - c) A 'lessons learned' summary (what you did do well, both technically and organizationally, and what you would do differently if you were to do it again). This summary should be brief (1-2 pages).

3.3 Due Dates and Mark Distribution

<i>Deliverable</i>	<i>Weight</i>	<i>Due Date</i> ⁴	<i>File Name</i>
RTX Project P1 Implementation ^{0,3}	10%	TBA ²	p1_Gid.ext
Software Design Description ^{0,1}	35%	TBA ²	sdd_Gid.ext
RTX Project P2 Implementation ^{0,3}	20%	TBA ²	p2_Gid.ext
RTX Project P2 Demonstration		TBA ²	
RTX Project P3 Implementation ^{0,3}	25%	TBA ²	p3_Gid.ext
RTX Project P3 Demonstration		TBA ^{2,4}	
Final Project Report ^{0,1}	10%	TBA ²	frpt_Gid.ext

Notes:

⁰ Only softcopy submission is sufficient. The name for RTX project P1 implementation, RTX project software design description, RTX project P2 implementation, RTX project P3 implementation and final project report should be *p1_Gid.ext*, *sdd_Gid.ext*, *p2_Gid.ext*, *p3_Gid.ext*, and *frpt_Gid.ext* respectively. Replace the 'id' with your two digit group number and "ext" with proper file extension. The ECE Course Book System is to be used for softcopy submission.

¹ Electronic submission in Microsoft Word format.

² To be confirmed with the class. The dates will be posted on the project web site.

³ Electronic submission. Put all source, header files, and binaries in a separate (submit) directory. Include a README file with group identification, description of directory contents and compilation procedure (and a makefile). Compress the directory contents into a single file. For archiving, you must choose one of *zip*, *tar.gz*, *tar.bz2* and *tar.Z*.

⁴ Groups that use part of or all of the 3 grace days for delayed submissions are not eligible for early project demonstration. A **5% of bonus** of the entire project will be applied to early demonstration groups.

3.4 Development and Demonstration Environment

Software development support includes *gcc*, a ColdFire/68K cross-compiler. The cross-compiler can be used to compile and link C and assembly code into a format that can be downloaded from a Nexus workstation into the MFC5307 based microcontroller. The cross-compiler is available on both Nexus and E&CE UNIX workstations.

The course laboratories are located in E2-2363. They contain Nexus workstations and MFC5307 based microcontroller boards. The demonstration of RTX operation (cf. Deliverables) will be done in the course laboratories. Project TAs and lab instructor will be available in their office hours to assist you with your implementation problems. Note that all the scheduled lab sessions for SE 350 are shown in Course Book System.

Additional Nexus workstations can be found at various locations on campus. Nexus workstations have X-terminal functionality providing remote access to LINUX/UNIX workstations.

You may prefer to do the initial development on a platform which is more convenient to access for you (e.g. a home computer). With some preparation, you could also test and debug most of your code on such a platform (you would need to use a simulator of the SBC I/O, see course web pages). However, as stated earlier, the project must be demonstrated on the SBC boards in the course laboratories.

4 Description of RTX Primitives & Services

This section lists the RTX primitive and services. You must implement these as described and may not modify the prototypes in any way. The primitives listed below will always return a value, either a pointer or an `int` return code. In the latter case, the return code value of 0 indicates success; non-zero value indicates a failure where applicable.

4.1 Interprocess Communication

The RTX will support a message-based IPC discussed in lectures. Messages are carried in envelopes (memory blocks, see below) with a header which is less than 64 bytes. Two IPC primitives will be implemented. The two primitives are:

```
int send_message (int process_ID, void * MessageEnvelope)
```

Delivers to the destination process a message carried in the message envelope (a memory block). Changes the state of destination process to `ready_to_execute` if appropriate. The sending process is preempted if the receiving process was blocked waiting for a message and has higher priority, otherwise the sender continues executing. The header of the message will have the layout given in the course overheads. It also fills in the *sender_process_id* and *destination_process_id* fields in the message envelope. The fields *sender_process_id*, *destination_process_id* and *message_type* are all of type `int`. The sender fills in the *message_type* field of the message envelope before invoking the primitive.

```
void * receive_message (int * sender_ID)
```

This is a blocking receive. If there is a message waiting, a pointer to the message envelope containing it will be returned to the caller. If there is no such message, the calling process blocks and another process is selected for execution. The sender of the message is identified through *sender_ID*, unless it is `NULL`.

4.2 Memory Management

The RTX supports a simple memory management scheme. The memory is divided into blocks of fixed size (128 bytes *minimum*). The size and the number of these blocks is a configuration parameter. The blocks can be used by the requesting processes for storing local variables or as envelopes for messages sent to other processes. A block which is no longer needed must be returned to the RTX. Two primitives are to be provided:

```
void * request_memory_block ()
```

The primitive returns a pointer to a memory block to the calling process. If no memory block is available, the calling process is blocked until a memory block becomes available. If several processes are waiting for a memory block and a block becomes available, the highest priority waiting process will get it.

```
int release_memory_block (void * MemoryBlock)
```

This primitive returns the memory block to the RTX. If there are processes waiting for a block, the block is given to the highest priority process, which is then unblocked. The caller of this primitive never blocks, but could be preempted. Thus, it may affect the currently executing process.

4.3 Processor Management

```
int release_processor ()
```

Control is transferred to the RTX (the calling process voluntarily releases the processor). The invoking process remains ready to execute. Another process may possibly be selected for execution.

4.4 Timing Services

```
int delayed_send (int process_ID, void * MessageEnvelope, int delay)
```

The invoking process does not block. The message (in the memory block pointed to by the second parameter) will be sent to the destination process (*process_ID*) after the expiration of the *delay* (timeout, given in msec units).

4.5 Process Priority

Process priorities have an integer *priority* value (0, 1, 2, 3, 4) where 0 is the highest priority level.

```
int set_process_priority (int process_ID, int priority)
```

This primitive sets the priority of the process with *process_ID* to the value given in *priority*. A process may change its own priority. The priority of the null process may not be changed from level 4 and it is the only process that can be assigned to level 4 (see also section 5.1). The caller of this primitive never blocks, but could be preempted. This preemption may affect the currently executing process.

```
int get_process_priority (int process_ID )
```

This primitive returns the priority of the process specified by the *process_ID* parameter. For an invalid *process_ID*, the primitive returns “-1”.

5 Required Processes

This section describes the processes which you must implement for the project.

5.1 System Processes

System processes are those processes needed by the system to perform basic services (scheduling and I/O).

5.1.1 Null Process

This process runs as the lowest priority process (level 4) in the RTX.. The Null process is the only process assigned to level 4. Level 4 is basically a “hidden” priority level reserved for the Null process. This preserves the 4 levels of user priorities (levels 0, .. , 3). Process id 0 is reserved for the null process. Initially, the following pseudocode can be used to design your null process:

```
loop forever
    release_processor( )
end loop
```

Once you have preemption working, then the “release_processor()” line could be removed from the infinite loop.

5.1.2 System Console I/O Processes

The system console is used for communication with the RTX and application processes. It consists of two devices: keyboard and CRT display. These two devices communicate serially with the microcomputer; using the receive and transmit lines of one of the two RS-232 ports.

The RTX will include two system processes, the Keyboard Command Decoder (KCD) process and the CRT Display process. These processes work in cooperation with the UART interrupt handler i-process.

5.1.2.1 Keyboard Command Decoder process

A keyboard command starts with the prompt character %, followed by a single (or multiple) letter command identifier and possibly additional command data. For example, %WS12:45:00 could be a command to the wall clock process, telling it to start the wall clock, setting the current time to 12:45:00 (where the command format is %WS hh:mm:ss).

The command decoder process responds to two types of messages: console keyboard input and command registration. The latter contains the command identifier and the process id of the process to which such commands are to be delivered when entered on the console keyboard. The processing of messages received depends on their type:

- 1. Command Registration**

The command identifier is associated with the registrant’s process id.

- 2. Keyboard Input**

The string input is sent to CRT display for output. If the string begins with a registered command identifier, it

is also sent to the registered requester.

5.1.2.2 CRT Display Process

This process responds to only one message type: a CRT display request. The message body contains the character string to be displayed. The string may contain control characters (e.g. newline). The process causes the string to be output to the console CRT. In printing to the console display, the process must use the UART i-process. Any message received is freed using the `release_memory_block` primitive.

5.2 Interrupt Processes (I-Processes)

Two interrupt handling processes are required:

5.2.1 Timer I-Process

The timer i-process is executed each time a hardware timer interrupt occurs. The timer i-process should handle the delivery of delayed send messages after the required time has expired.

5.2.2 UART I-Process

The UART i-process uses interrupts for both the transmission and receiving of characters from the serial port. No polling or busy waiting strategies may be implemented. The UART i-process forwards characters (or commands) received to the KCD, and also responds to messages received from the CRT display process to transmit characters to the serial port.

5.2.2.1 Hot Keys

As well, the UART i-process is used to provide debugging services which will be used during the demonstration. Upon receiving specific characters (*hot keys* - your choice, e.g., !) as input, the UART i-process will print the following to the debug console:

1. The processes currently on the ready queue(s) and their priority.
2. The processes currently on the blocked on memory queue(s) and their priorities.

As well, you are free to implement other hot keys to help in debugging. For example, a hot key which lists the processes, their priorities, their states; or another which prints out the number of memory blocks available. Like all other debug prints, the hot key implementation should be wrapped in

```
#ifdef _DEBUG_HOTKEYS
...
#endif
```

preprocessor statements and should be turned off during automated testing. If the automated test processes fail, you may be asked to turn the hot keys on again in determining why the test processes are failing.

Another hotkey debug printout may be used to display recent interprocess message passing. A (circular) log buffer keeps track of the 10 most recent `send_message` and `receive_message` invocations made by the processes; upon receiving a specific hotkey, these most recent 10 sent and 10 received messages are printed to the debug console. The number 10 is used only as an example. The information printed could contain information such as:

1. Sender process id
2. Destination process id
3. Message type
4. First 16 bytes of the message text
5. The time stamp of the transaction (using the RTX clock)

5.3 User Processes

These processes will be used to demonstrate the operation of your system.

5.3.1 Set Priority Command Process

This process registers itself with the Keyboard Command Decoder process as the handler for the %C command. The %C command has two parameters:

```
% C process_id new_priority
```

where *process_id* and *new_priority* are integers. This command changes the priority of the specified process, *process_id*, to *new_priority*. The change in priority level is immediate. It could also affect the target process's position on a ready queue or a blocked resource queue. The parameters must be verified to ensure a valid *process_id* and priority level is given. A %C command with illegal parameters will be ignored with an error message printed on the console.

5.3.2 Wall Clock Display Process

This process registers itself with the Keyboard Command Decoder process as the handler for the %W command. The %WS hh:mm:ss command sets the current wall clock time to hh:mm:ss, starts the clock running and causes display of the current wall clock time on the console CRT. The display will be updated every second. The %WT command will cause the wall clock display to be terminated.

5.3.3 Demonstration of User-Level Processes

An important category of software tests are the stress tests. These tests seek to verify the behavior of the system under heavy stress scenarios. One such scenario is depletion (or near depletion) of system resources. For the demonstration of this project, you will implement three processes whose behavior is described below. The stress scenario being tested is depletion of memory blocks.

Process A:

```
register with Command Decoder as handler of %Z commands
loop forever
  receive a message
  if the message contains the %Z command then exit the loop
endloop

num = 0
loop forever
  request memory block to be used as a message envelope
  set message_type field to "count_report"
  set msg_data[0] field to num
  send the message to process B

  num = num + 1
  release_processor()
endloop
// note that Process A does not deallocate any received envelopes!
```

Process B:

```
loop forever
  receive a message
  send the message to process C
endloop
```

Process C:

```
perform any needed initialization and create a local message queue
loop forever
  if (local message queue is empty)
  then receive a message
  else dequeue the first message from the local message queue
  if msg_type == "count_report"
```

```

    then if msg_data[0] is evenly divisible by 20
        then send "Process C" to CRT display using msg envelope
            hibernate for 10 sec
        endif
    endif
    deallocate message envelope
    release_processor()
endloop

```

The line “**hibernate for 10 sec**” is further expanded as:

```

    request a delayed_send for 10 sec delay with msg_type=wakeup10
    loop forever
        receive a message          //block and let other processes execute
        if (message_type == wakeup10)
            then exit this loop
        else put message on the local message queue for later processing
    endloop

```

Notes:

- Process C has a local message queue (distinct from the incoming message queue maintained by the RTX) onto which it enqueues (in FIFO order) messages which arrive while it hibernates. It processes these messages later.
- For your own testing, set the priority levels for processes A, B and C to values which are most likely to cause memory block depletion in the RTX. During project demo, you may be asked to re-initialize your RTX with TA/instructor specified priorities for A, B, and C and vary the total number of message envelopes available.
- It is recommended that test processes A, B and C have process ids 7, 8, and 9 respectively. If you choose not to do this, you should have this information ready before the demonstration begins.

5.3.4 User-Level Test Process

Write two user-level test processes to test your own OS and another project group's OS. These two test processes should run in user mode and do not assume any kernel level data structures. These two test processes only call the RTX API. A header file will be provided later during the term to provide the interface of set and get a user-level test process's process id, initial priority, stack size and entry point. The test processes should finish testing within 5 seconds.

6 RTX Initialization

To make the RTX more generally applicable, the RTX will be configured at initialization as specified in the RTX Configuration Table. This table has three sections:

1. Memory configuration section: memory block size, number of memory blocks created;
2. System process section;
3. Application process section.

The process sections list the processes to be created. Each entry contains the following data: process id, priority, stack size, start address, and for system processes, whether the process is an i-process. You will also be given a header file for the initialization of automated test processes. *All initializations must take place after the RTX execution starts.*

7 Automated Testing

The demonstration of your working project will consist of two phases. In the first phase, initial manual testing will be performed using the specified keyboard commands. There should be exactly 30 memory blocks available to the user processes for this part of the demonstration.

In the second phase, you will be required to link 6 test processes (supplied at time of demo) with your RTX and execute the automated test sequence. To ensure that you can successfully link with the test processes, you will be provided with a test related header file and several sets of six dummy test processes. The header file contains all the necessary process related information needed during system initialization. The first set will contain processes each of which are represented by:

```

loop forever
    release_processor( )
end loop

```

A second set is provided as object files which display simple messages and are provided to ensure that your compilation script can accommodate them. Some minimal tests are provided, but you should not rely on them to ensure that your system is fully functional.

One requirement is that the six automated testing processes have reserved process ids 1, 2, 3, 4, 5 and 6.

Since the test processes have no knowledge of your detailed internal design, they only invoke the functions specified by the RTX API. In addition, to allow you freedom in your design, the test processes only assume that your message envelope format conforms to the general guidelines given in lectures. The basic restriction is that your message envelope be at least 128-bytes in length. Test processes will use an offset of 64-bytes from the pointer (to a message envelope) and use that area (within the message text, see Fig. 1) for inter-test process communication. So, none of your primitives should modify this message text area while handling a message envelope (i.e. deallocate, allocate, delay, etc.). This area, within the message text, is only used by the test processes while they “own” an envelope. Your own processes can do anything they want with this area when they “own” an envelope.

The test processes use the MCF5307’s Timer 1 allowing your software to use Timer 0. The UC and the UD bits in the MBAR (Module Base Address Register) should be left unset. The test processes use the microcomputer board’s debug monitor to print messages and expect that your RTX does not use the debug monitor while the test processes are executing.

It is recommended that all debug statements be wrapped in `#ifdef DEBUG ... #endif` preprocessor statements so that they may be turned off as a group during the automated testing. Turning these off may also affect timing performance.

An overview of the expected mapping of process ids to processes is given in the table below:

Process_id	Owner Process
0	Null process
1	Test_process_1
2	Test_process_2
3	Test_process_3
4	Test_process_4
5	Test_process_5
6	Test_process_6
7	Process_A
8	Process_B
9	Process_C
10 ...	Other processes start at 10

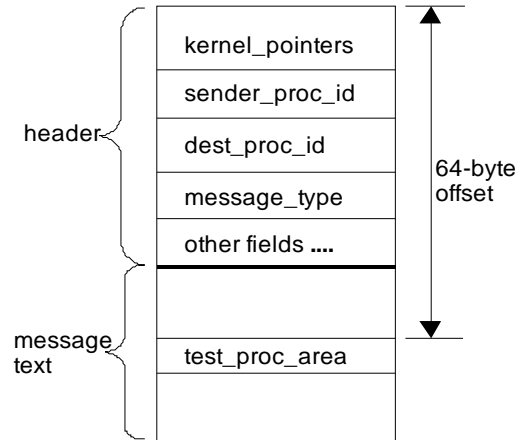


Fig. 1: Typical Message Envelope Format

Finally, you will be asked to link your RTX with test processes written by other project groups. Each project group should write two user-level test processes as specified in section 5.3.4 for other project groups to use.