

RTOS Functional Overview

- Process Management
 - create/terminate processes
 - perform reset or 'power-on' initialization
- Processor Scheduling
 - selection & dispatch of processes for execution
- Inter-Process Communication/Synchronization
 - send/receive messages
 - signal/wait on semaphores
- Storage Management
 - allocate/deallocate memory

RTOS Functional Overview [2]

- Interrupt Handling Framework
 - capture interrupts
 - if required, activate a user defined process
- Timing Services
 - relative [delay] services
 - absolute time services
- Device Driver Interfaces
 - provide standard i/o and specialized interrupt driven device handlers

Simple RTOS: Requirements

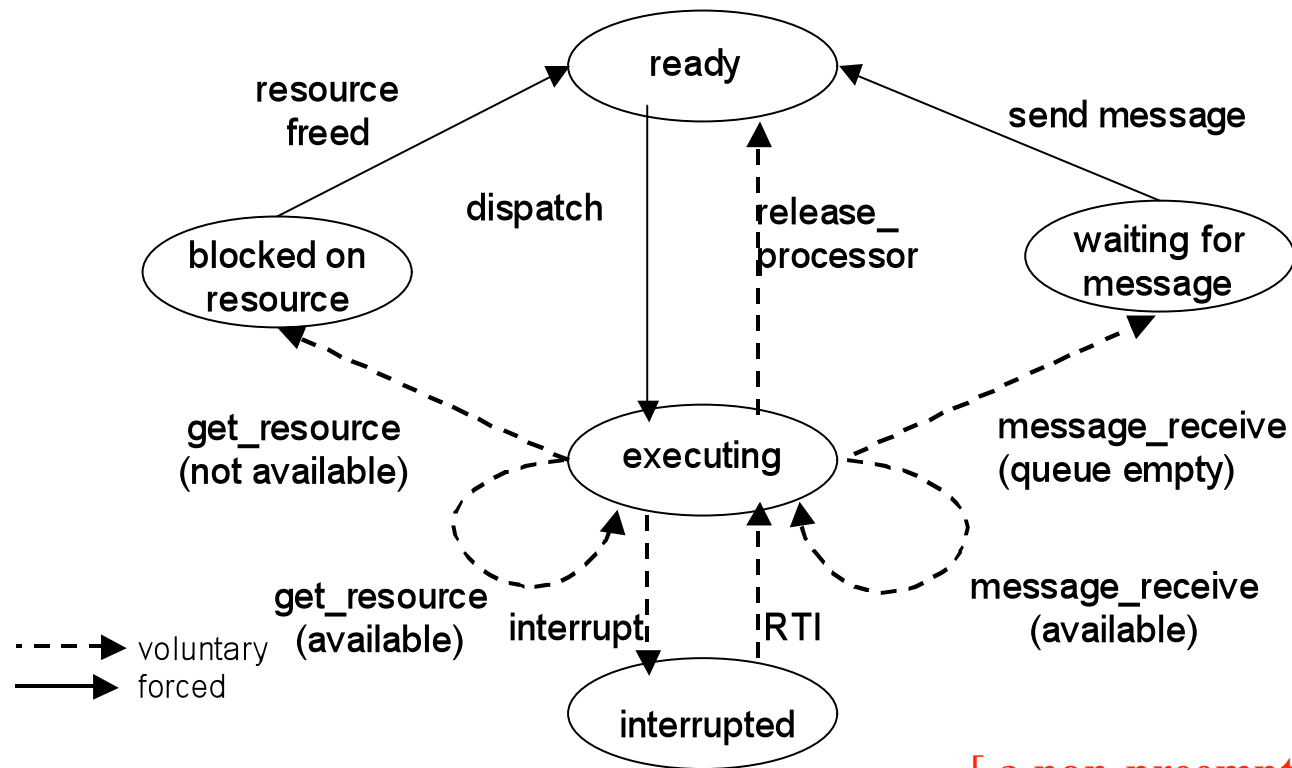
- basic requirements
 - non-preemptive
 - support for processes [creation at init time only]
 - priority scheduling [fixed priority]
 - message-based interprocess communication [asynchronous, messages sent in envelopes]
 - memory management: message envelopes
 - basic timing services

Simple RTOS: Setting

- Assumptions concerning RTOS setting
 - all processes are known and created at OS initialization time
 - processes are friendly, cooperating and non malicious
 - each process 'knows' the process_id of its co-workers

RTOS Process States

- simplified



[a non-preemptive system]

RTOS: Atomicity

[conceptual only]

- RTOS primitives must execute indivisibly
- define private kernel function *atomic(on / off)*
 - *atomic(on)* enables the atomic functionality
 - first executable statement in each primitive
 - *atomic(off)* disables the atomic functionality
 - last statement in each primitive (before 'RET')
- possible implementation
 - extra field: `int atomic_count = 0;`
 - *atomic(on)* increments, *atomic(off)* decrements this field
 - whenever `atomic_count > 0`, atomicity must be enforced by kernel

in each PCB

RTOS: Atomicity Cont'd

- if there is direct access to CPU interrupt masking, then explicit *atomic()* function is not needed
- possible implementation of *atomic*
 - *on*: save interrupt system mask, mask all interrupts
 - *off*: restore interrupt system mask

must use assembly language instructions


RTOS: Atomicity Cont'd

- in the following slides, the *atomic(on/off)* functionality is used as to indicate the need to enforce indivisibility of kernel primitives
- Question:
 - must the *atomic(on/off)* functionality be implemented exactly as described?
- Answer:
 - discussed in lectures.....

Atomic(on/off) concept used only in slides to stress the requirement of atomicity -- since in reality, all the primitive functionality will be performed in “kernel mode” with interrupts disabled by a kernel version of the user visible API functions, after the user API ‘traps’ into the kernel.

RTOS: *current_process*

- RTOS must know which process **currently executes**
- RTOS design includes a private kernel variable *current_process*
- *current_process* always refers to the currently executing process (more exactly, to its internal representation, e.g. process object or PCB data structure)



RTOS: *process_switch()*

- frequently needed procedure: remove the currently executing process from the CPU, select the next process to execute and give the CPU to it
- RTOS design includes a private kernel function *process_switch()*
 - invokes the scheduler to select the next process to be executed
 - invokes *context_switch(next_process)*

RTOS: *context_switch(next_proc)*

- save context of currently executing process into its PCB/process object
- sets *current_process* to refer to *next_proc*
- sets the state of the *current_process* to executing
- restores the context of *current_process*
- causes the *current_process* execution to begin

RTOS: *process_switch()* cont'd

- after a primitive executes *process_switch*:
 - the invoking process will eventually regain control again
 - execution will eventually resume on the instruction immediately following the *process_switch* instruction
 - when?

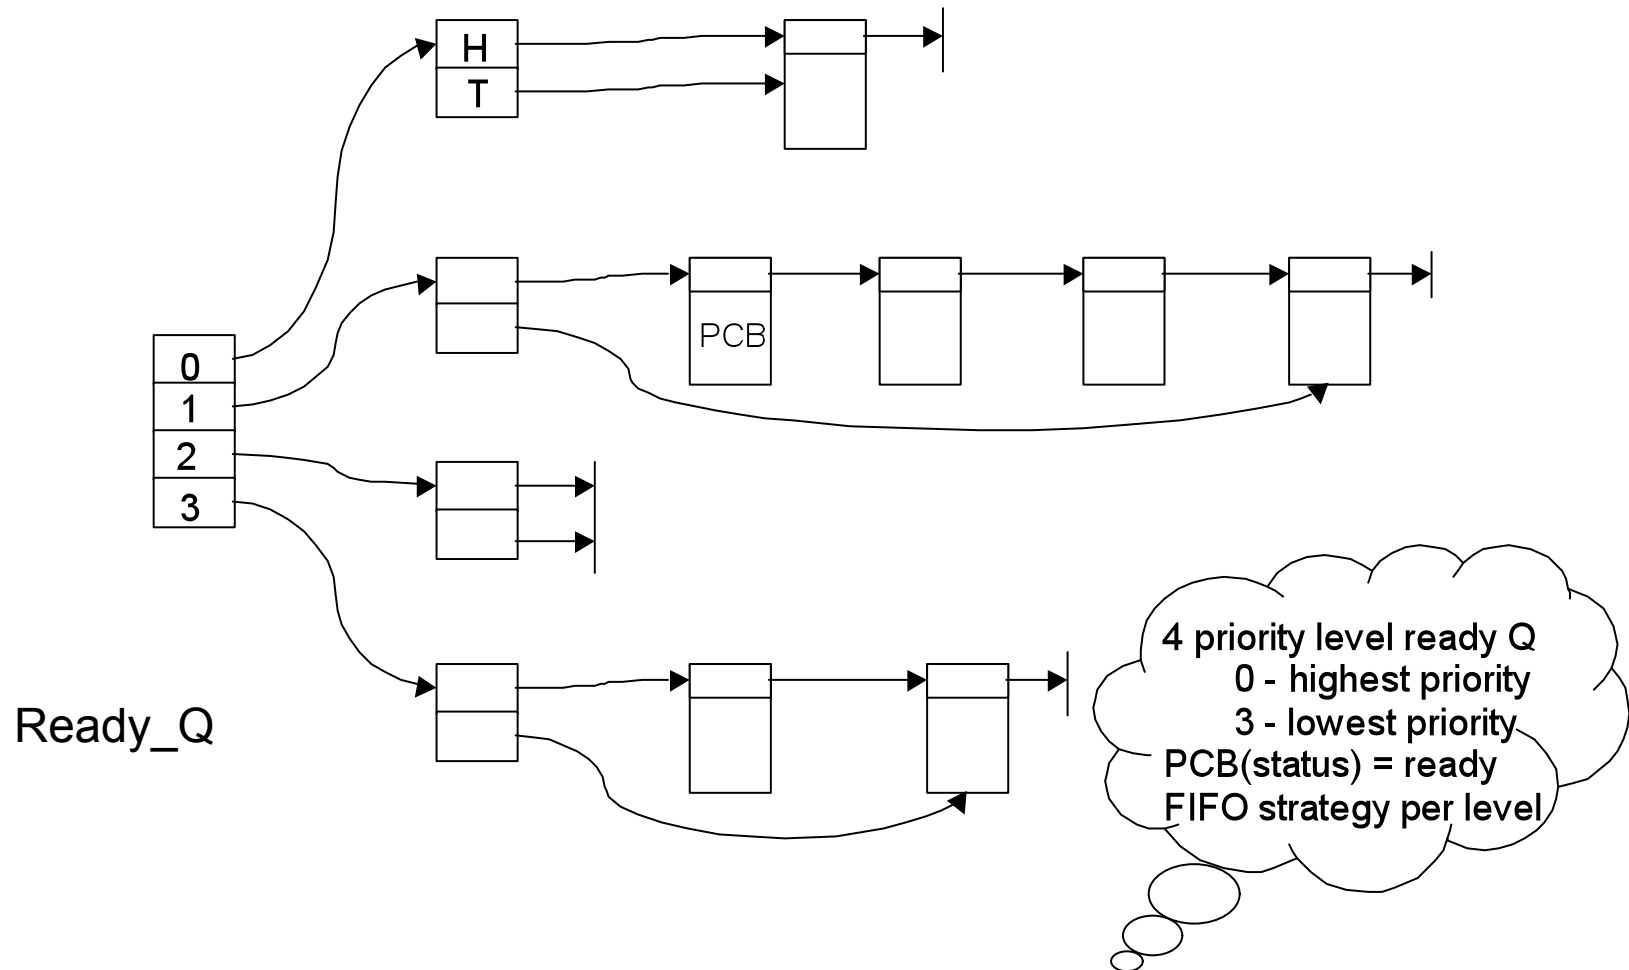
RTOS: Scheduling

- requirement spec: fixed priority based
- each process assigned a priority (urgency)
 - highest priority ready process will get to execute
 - processes with equal priority treated as FCFS
- applicable to both preemptive and nonpreemptive RTOS
- possibility of indefinite-blocking (starvation, livelock)
 - arrival rates of high priority processes may be so high that a low priority process may wait for extended time/forever to execute

RTOS: Scheduling

- *process_switch()* invokes the scheduler
- scheduler
 - selects the highest priority ready process
- *process_switch* then invokes *context_switch(next_proc)* to let the selected process execute
- note: other scheduling algorithms used in more complex real-time operating systems

RT OS: Priority Scheduling



Real-Time OS – *rpq_enqueue/dequeue*

- fixed priority based scheduling \Rightarrow
our design includes:
 - private kernel functions
 - rpq_enqueue(PCB/proc object)*
enqueues the PCB/process object on the
appropriate ready process queue based on its
priority
 - rpq_dequeue()*
dequeues and returns reference to **highest-
priority ready process**

RTOS: Null Process

- CPU must always execute something
- what should the RTOS do when the scheduler finds that the ready queue is empty?
 - possible solution:
 - loop within RTOS, periodically check **(be careful!)**
 - make sure that the ready queue is never empty!!
- how?
 - include a process (*null process*) with the lowest priority that is always ready to run

RTOS: Null Process Cont'd

- basic null process functionality:

```
    null_process:  
        while (true) {  
            release_processor();  
        }
```

- should the null process do more than that?
- two views
 - strict view: one process, one function, hence no
 - permissive view: let it do something useful
 - e.g. ROM checksum check, low level OS checks

RTOS: *release_processor()*

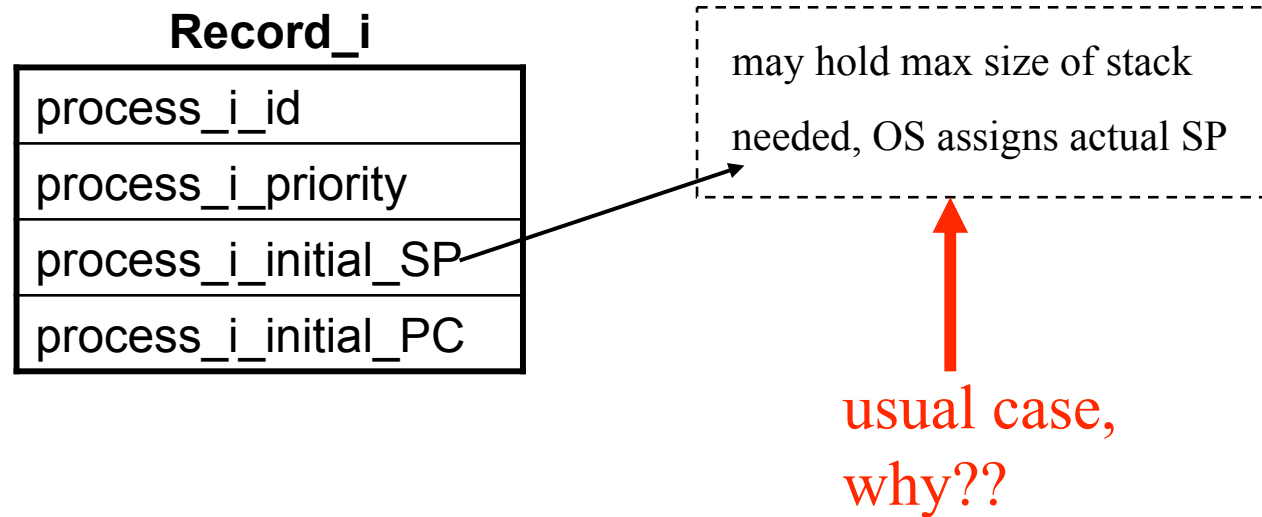
- RTOS design includes the following primitive:
release_processor()
- *release_processor*:
 - set *current_process* state to ready
 - *rpq_enqueue (current_process)*
 - *process_switch()*

RTOS: Initialization

- what operations need to be carried out at OS startup (i.e. after power up, reset)?
- initialize all HW, OS structures, create processes and start process execution

RTOS: Process Initialization Table

- RTOS must know which processes to create
- our design: array of records (initialization table, IT)
- each record contains the information necessary to start its respective process
- a record could have the following structure:



RTOS: Initialization Sequence

- during initialization, RTOS
 - initializes all hardware
 - creates and initializes all kernel data structures
 - reads IT
 - creates PCBs/process objects as needed (proc_status=ready)
 - places each PCB into its respective scheduling ready queue
 - invokes scheduler to select first process to execute
 - lets the selected process start executing

RTOS: IPC

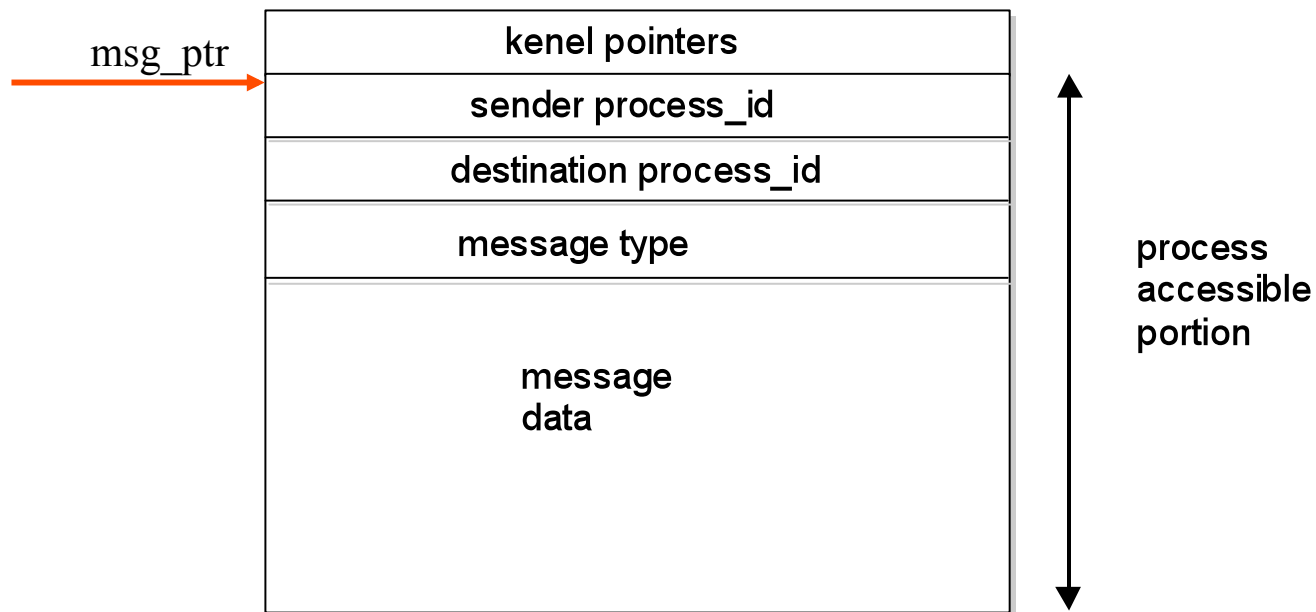
- requirement spec:
 - message-based, asynchronous IPC
 - messages carried in shared message blocks (msg envelopes)
- each process writes a message into a msg envelope
- process invokes *send(msg_envelope)*
- issues:
 - what is the format of the message envelope (i.e. envelope memory block)?
 - where do these memory blocks come from?

RTOS: Message Envelopes

- message envelopes managed by the kernel
 - an appropriate number of message envelopes (blocks of memory) created at system init time
 - a process allocates a message envelope to send a message
 - a process deallocates an envelope when it is no longer needed (current owner of the envelope!)
 - a process owns a message envelope that it receives or allocates (until it is sent)


RTOS: Msg Envelope Format

- design of message envelope format:
 - fixed size block of memory
 - layout:



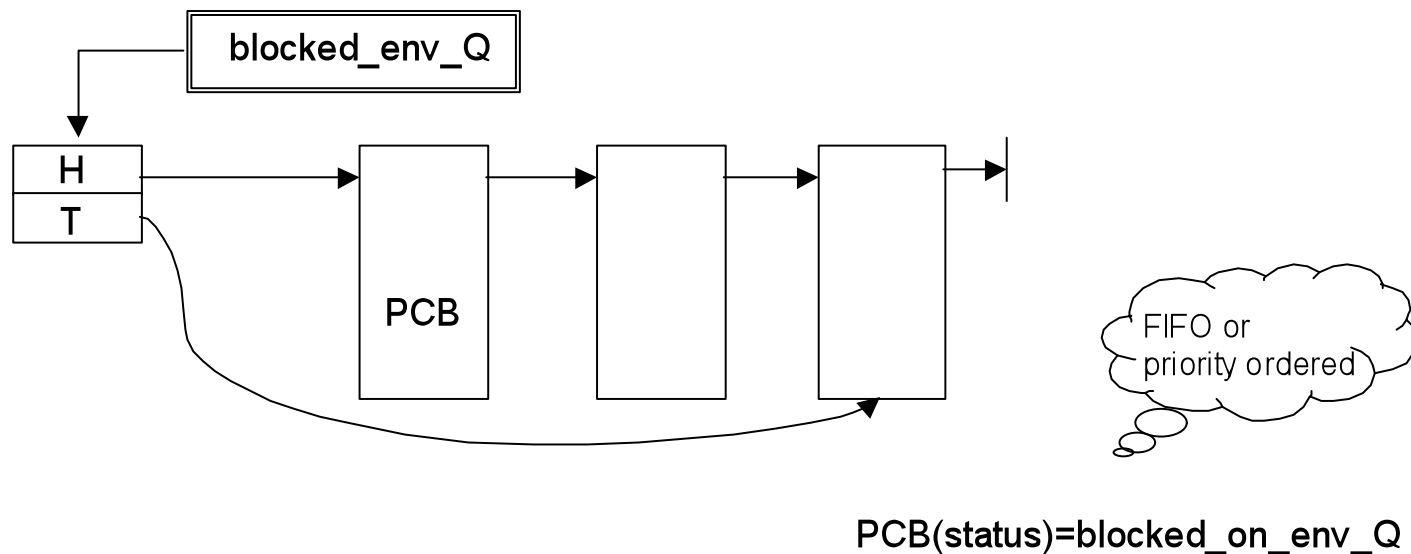
RTOS: Msg Envelope Management

- where do message envelopes come from?
 - at initialization, kernel creates a fixed number of envelopes and put them on a free envelope queue



RTOS: Msg Envelope Management

- if a process does not already have an envelope then it first allocates an envelope from the kernel
- what if no envelopes left?
 - requesting process blocks
 - our design: blocked processes kept on blocked_env_Q



RTOS: *allocate_envelope*

- functionality of *allocate_envelope* primitive:

```
allocate_envelope() {  
    atomic(on);  
    while (free_env_Q is empty) {  
        put process object/PCB on blocked_env_Q  
        set process state to blocked_on_env_allocate  
        process_switch();  
        ***restart here when blocked process executes eventually  
    }  
    env ← reference to de-queued envelope  
    return env;  
    atomic(off);  
}
```

RTOS: *deallocate_envelope*

- functionality of *deallocate_envelope*:

```
deallocate_envelope( in: env ) : {  
    atomic(on);  
    put env onto free_env_Q  
    if ( blocked_env_Q not empty)  
        { dequeue one of the blocked processes  
          set its state to ready and enqueue it on ready process queue }  
    atomic(off);  
}
```




- invoking process never blocks!
- how does it work with respect to *allocate_envelope* kernel primitive??

RTOS: IPC

- how does a process send/receive a message?
- requirement spec: message-based, asynchronous
- design decision
 - non-blocking *send*
 - blocking *receive*
- design issue:
 - how are messages buffered by kernel?
 - if multiple processes send a message to a process but that process does not do a receive for some time
 - how does kernel keep track of such messages?

RTOS: Waiting Messages

- design: let each process have a queue of waiting messages
 - extend the PCB to include:



RTOS: IPC

- what happens to a process that executes *receive* but no message available?
 - it blocks
 - its state is set to `blocked_on_receive`
- design issue: should processes in this state be kept somewhere (queue, set ?)
 - why put on some Q?
 - no real need to do so.
 - just set its status to “`blocked_on_receive`” and that’s all

RTOS: *receive*

- functionality of *receive* primitive:

```
receive() {  
    atomic(on);  
    while ( current_process's msg_queue is empty) {  
        set state of current_process to blocked_on_receive  
        process_switch( );  
        *** return here when this process executes again  
    }  
    env ← dequeued envelope from the process' message queue  
    return env  
    atomic(off);  
}
```

RTOS: *send*

- functionality of *send* primitive:

```
send( target_pid, env) : {  
    atomic(on);  
    set sender_procid, destination_procid fields in env  
    target_proc ← convert target_pid to process obj/PCB ref  
    enqueue env onto the msg_queue of target_proc  
    if ( target_proc.state is blocked_on_receive)  
        { set target_proc state to ready  
          rpq_enqueue( target_proc );  
        }  
    atomic(off);  
}
```

What about a preemptive system?

RTOS: Interrupt Handling

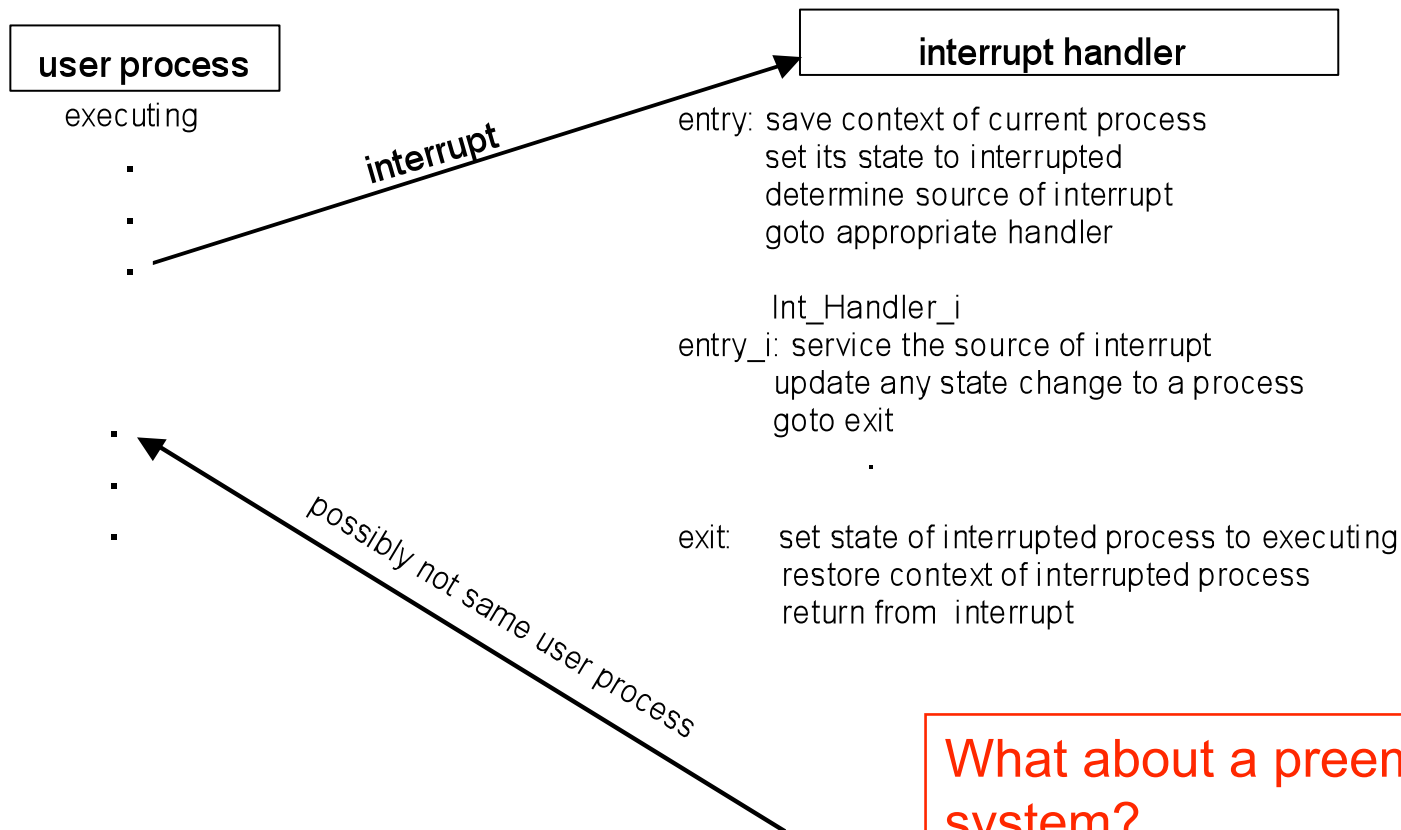
- in real-time OS, interrupt handling must be fast
 - short latency to respond to interrupt
 - fast processing by interrupt handler
- interrupts may cause a change in state for some blocked process
 - e.g. a process blocked for external event to occur will have its state changed to ready and be placed on the ready process queue when the event occurs

RTOS: Interrupt Handling Issues

- possible interpretation of an interrupt:
 - an interrupt is a *hardware message* usually requiring a short latency and quick service
- design issues:
 - does the interrupt handling code run as part of kernel, within a process (if yes, which?)
 - are interrupt handlers themselves interruptible?
 - if OS is preemptive, need to deal with the possibility that an interrupt results in a higher priority process becoming ready

RTOS: Interrupt Handling Sequence

- abstracted interrupt processing sequence (nonpreemptive system)




What about a preemptive system?

RT OS: Interrupt Handler Design

- interrupt handler must interact with OS processes
- alternatives:
 - multiple ad-hoc interaction mechanisms
 - i_process
 - an i_process gets the CPU from an interrupt handling sequence, not through the dispatcher.
 - never blocks if it invokes a kernel primitive
 - the interrupt (exception) handling routine ~~and~~ starts the appropriate i-process.
 - conceptually: i-process has the max priority, is scheduled by interrupt

RTOS: i-process

- state diagram for an i_process:



- a process object/PCB is associated with each i_process with state = i_process (permanently)
- always ready to run, but not on any ready Q

RTOS: i-process constraints

- an i_process can invoke kernel primitives:
 - however, an i_process is not allowed to block!
- primitives which can block a process **must** be modified to ensure that i_process does not block!
 - e.g. the synchronous receive message primitive
 - return null if the invoking process is an i_process and there is no message waiting
 - similarly for other primitives

RTOS: Interrupt Handling

- we can now detail the previous exception handling sequence
- note: *save_proc* stores reference to the process/PCB which was executing (was the *current_proc*) when the interrupt occurred

exception_handler:

begin

set the state of current process to interrupted

save_proc = *current_process*

select interrupt source

A: *context_switch* (*i_proc_A*)

break

RTOS: Interrupt Handling

```
Z:          context_switch (i_proc_Z);  
    break  
end select  
//code to save context of interrupt handler (i_process)  
current_process = save_proc;  
context_switch (save_proc);  
//perform a return from exception sequence  
//this restarts the original process before i_handler  
end;
```

RTOS: Timing Services

- fundamental service in real-time operating systems
- service categories:
 - sleep: defer execution for n seconds
 - voluntarily give up CPU until the specified time expires, then be put back on ready queue
 - timeout notification:
 - request kernel to inform process when a specified time period has expired; process continues execution
 - repetitive timeout notification:
 - repeated timeout notification until cancelled

RTOS: Timing Services 2

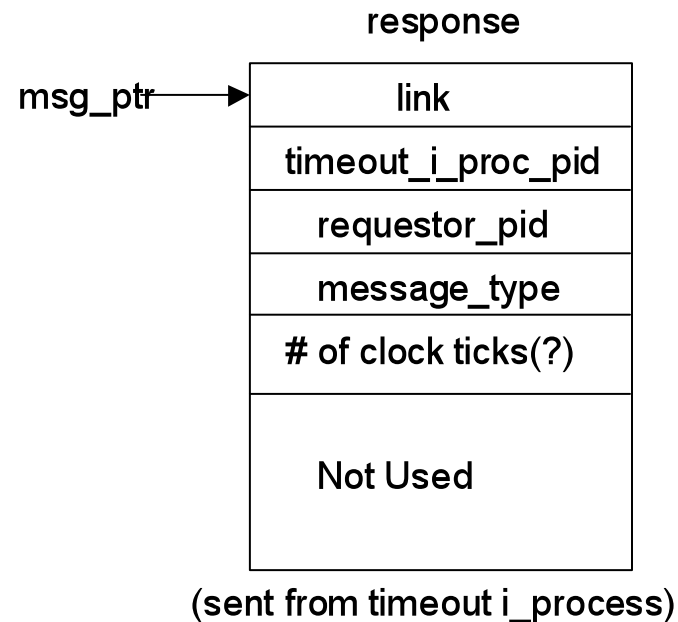
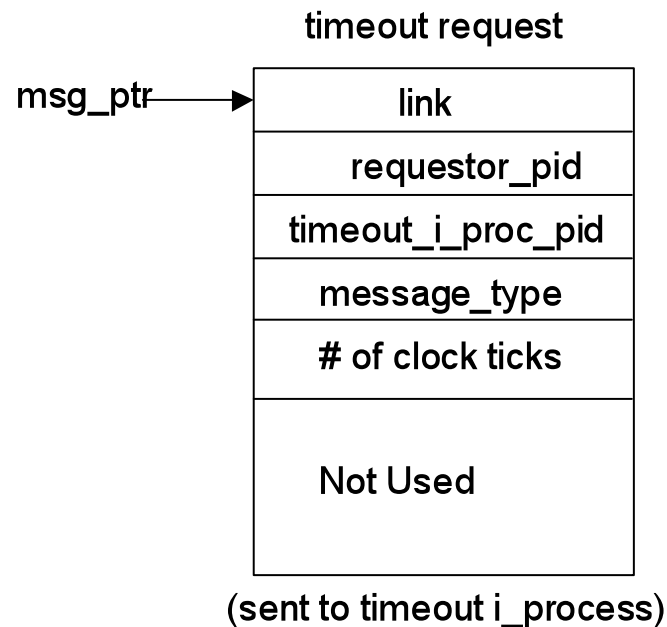
- service requests could be stated in:
 - relative time (x clock ticks)
 - absolute time (February 24, 10:34:22 AM, 2002)
 - others as appropriate
- related functionality
 - cancellation of earlier request

RTOS: Timing Service Design

- two parts: interface/protocol design, internal design
- interface/protocol design
 - basic service only (timeout), no cancellation
 - service request, expiry notification: by messages
- internal design
 - timing service implemented by *i_process*
 - service request: a user process send a request message to the timing *i_process*
 - timeout notification: the *i_process* sends a message back

RTOS: Timing Service Request


- *send(timeout_i_process, message)*
 - message contains a timeout request
 - request format:



RTOS: Timing Service Request 2

- user processes 'know' the pid of the timeout i_process
- after the expiration of the time, the timeout i_process sends the original message envelope back to requestor
- timeout service maintains requests in a sorted list

RTOS: Timing Service Design



- to reduce CPU overhead, expiry_time can be replaced by the # of clock ticks after the expiry of the predecessor in the list
- example: queue timeout list {25, 30, 0, 10}
 - one timeout for 25 clock ticks
 - two timeouts for 55 clock ticks
 - one timeout for 65 clock ticks

RTOS: timeout i-process

- at each clock tick (hardware timer interrupt), the timeout i-process:
 - increments `current_time`
 - invokes *receive()* to see new requests
 - since it is an `i_process`, it cannot block!!
 - if any new requests, adds them to sorted list
 - checks whether any timing requests have expired, if yes, it sends the notification using the request message envelope back to the requester

RTOS: timeout i-process 2

- basic outline:

timeout_i_process:

{

env \leftarrow receive(); //to get pending requests

while (env is not null)

{

//code to insert the envelope into the timeout_list

env \leftarrow receive(); //see if any more requests

}

//continued next slide

RTOS: timeout i-process 3

```
if (timeout_list is not empty)
{
    while (head_of_timeout_list.expiry_time <=
cur_time)
    {
        env ← timeout_list.dequeue();
        target_pid ← env.source_pid;
        env.source_pid ← timeout_i_process_pid;
        send( target_pid, env ); //return envelope
    }
}
}
```

RTOS: Preemption

- certain RTOS primitives can make ready a process whose priority is higher than that of current process
 - *send*
 - *deallocate_envelope*
- pre-emptive OS \Rightarrow highest priority ready process should execute
- preemption is relatively easy to add to the design presented

Design Changes for Preemption

- on return from *send*, etc primitives, include check
 - if hpr process priority $>$ current process priority, do context switch
- problem when i-process executes
- possible solution
 - let priority of i-processes be max priority
 - on return from interrupt, check whether the priority of the interrupted process still the highest
 - if not, context switch to hpr process

entry

